

model

```

In [1]: class NeuralNetMLP(object):
    def __init__(self, is_sigmoid=True, initializer=None, dropout_rate=1., n_layers=
        n_hidden=5, l2=0., epochs=100, eta=0.0005, shuffle=True,\
            minibatch_size=64, seed=666):
        self.random = np.random.RandomState(seed)
        self.is_sigmoid = is_sigmoid
        self.initializer = initializer
        self.dropout_rate = dropout_rate
        self.n_layers = n_layers
        self.n_hidden = n_hidden
        self.l2 = l2
        self.epochs = epochs
        self.eta = eta
        self.shuffle = shuffle
        self.minibatch_size = minibatch_size
        self.k = locals()
        if n_layers < 3:
            raise Exception("n_layers must be >= 3 !!!")
    def _onehot(self, y, n_classes):
        """Encode labels into one-hot representation
        Parameters
        -----
        y : array, shape = [n_samples]
        Target values.
        Returns
        -----
        onehot : array, shape = (n_samples, n_labels)
        """
        onehot = np.zeros((n_classes, y.shape[0]))
        for idx, val in enumerate(y.astype(int)):
            onehot[val, idx] = 1.
        return onehot.T
    def _sigmoid(self, z):
        """Compute logistic function (sigmoid)"""
        return 1. / (1. + np.exp(-np.clip(z, -250, 250)))
    def _forward(self, X, is_training=True):
        """Compute forward propagation step"""
        # step 1: net input of hidden layer
        if is_training:
            dropout_rate = self.dropout_rate
        else:
            dropout_rate = 1.
        for i in range(self.n_layers - 2):
            if i == 0:
                # [n_samples, n_features] dot [n_features, n_hidden]
                # -> [n_samples, n_hidden]
                r = self.random.binomial(1, dropout_rate, size=X.shape)
                X = r * X
                self.k['z_h_{0}'.format(i+1)] = \
                    np.dot(X, self.k['w_h_{0}'.format(i+1)]) + self.k['b_h_{0}'.format(i+1)]
                # step 2: activation of hidden layer
                if self.is_sigmoid:
                    self.k['a_h_{0}'.format(i+1)] = self._sigmoid(self.k['z_h_{0}'.format(i+1)])
                else:
                    self.k['a_h_{0}'.format(i+1)] = self.k['z_h_{0}'.format(i+1)]
            else:

```

```

# [n_samples, n_hidden] dot [n_hidden, n_hidden]
# -> [n_samples, n_hidden]
r = self.random.binomial(1, dropout_rate, size=self.k['a_h_{}'.format(i)].shape)
self.k['a_h_{}'.format(i)] = r * self.k['a_h_{}'.format(i)]
self.k['z_h_{}'.format(i+1)] = \
    np.dot(self.k['a_h_{}'.format(i)], self.k['w_h_{}'.format(i+1)])
    + self.k['b_h_{}'.format(i+1)]
# step 2: activation of hidden layer
if self.is_sigmoid:
    self.k['a_h_{}'.format(i+1)] = self._sigmoid(self.k['z_h_{}'.format(i+1)])
else:
    self.k['a_h_{}'.format(i+1)] = self.k['z_h_{}'.format(i+1)]
# step 3: net input of output layer
# [n_samples, n_hidden] dot [n_hidden, n_classlabels]
# -> [n_samples, n_classlabels]
z_out = np.dot(self.k['a_h_{}'.format(self.n_layers - 2)], self.w_out) + self.b_out
# step 4: activation output layer
if self.regression_problem:
    a_out = z_out
else:
    a_out = self._sigmoid(z_out)
return [self.k['z_h_{}'.format(i+1)] for i in range(self.n_layers - 2)], \
        [self.k['a_h_{}'.format(i+1)] for i in range(self.n_layers - 2)], z_out, a_out

def _compute_cost(self, y_enc, output):
    """Compute cost function.
    Parameters
    -----
    y_enc : array, shape = (n_samples, n_labels)
    one-hot encoded class labels.
    output : array, shape = [n_samples, n_output_units]
    Activation of the output layer (forward propagation)
    Returns
    -----
    cost : float
    Regularized cost
    """
    L2_term = (self.l2 *
               (np.sum([np.sum(self.k['w_h_{}'.format(i+1)] ** 2.) for i in range(self.n_layers - 2)] +
                       [np.sum(self.w_out ** 2.)]))
               )
    term1 = -y_enc * (np.log(output))
    term2 = (1. - y_enc) * np.log(1. - output)
    if not self.regression_problem:
        cost = np.sum(term1 - term2) + L2_term
    else:
        cost = np.sum((output - y_enc) ** 2. )
    return cost

def predict(self, X):
    """Predict class labels
    Parameters
    -----
    X : array, shape = [n_samples, n_features]
    Input layer with original features.
    Returns:
    -----
    y_pred : array, shape = [n_samples]
    Predicted class labels.
    """

```

```

z_h, a_h, z_out, a_out = self._forward(X, is_training=False)
if self.regression_problem:
    y_pred = a_out
else:
    y_pred = np.argmax(z_out, axis=1)
return y_pred
def R_square(self, output, targets):
#     n = len(targets)
#     MSE = (1 / n) * (np.sum((output - targets) ** 2))
#     rMSE = MSE / np.var(targets)
#     r_2 = 1 - rMSE
return r2_score(targets, output)

def fit(self, X_train, y_train, X_valid, y_valid):
    """ Learn weights from training data.
    Parameters
    -----
    X_train : array, shape = [n_samples, n_features]
    Input layer with original features.
    y_train : array, shape = [n_samples]
    Target class labels.
    X_valid : array, shape = [n_samples, n_features]
    Sample features for validation during training
    y_valid : array, shape = [n_samples]
    Sample labels for validation during training
    Returns:
    -----
    self
    """
    if np.sum(y_train % 1) == 0:
        n_output = np.unique(y_train).shape[0] # no. of class labels
        self.regression_problem = False
    else:
        n_output = 1
        self.regression_problem = True
    n_features = X_train.shape[1]
    #####
    # Weight initialization
    #####

    for i in range(self.n_layers - 2):
        # weights for input -> hidden
        self.k['b_h_{}'.format(i+1)] = np.zeros(self.n_hidden)
        if i == 0:
            if self.initializer == 'Xavier':
                self.k['w_h_{}'.format(i+1)] = \
                    self.random.normal(loc=0.0, \
                                        scale=(2 / (n_features + self.n_hidden)) *
                                        , size=(n_features, self.n_hidden))
            else:
                self.k['w_h_{}'.format(i+1)] = \
                    self.random.normal(loc=0.0, scale=0.1\
                                        , size=(n_features, self.n_hidden))
        else:
            if self.initializer == 'Xavier':

```

```

        self.k['w_h_{}'.format(i+1)] = \
        self.random.normal(loc=0.0, \
                             scale=(2 / (self.n_hidden + self.n_hidden)
                                     , size=(self.n_hidden, self.n_hidden))

    else:
        self.k['w_h_{}'.format(i+1)] = \
        self.random.normal(loc=0.0, scale=0.1\
                             , size=(self.n_hidden, self.n_hidden))

# weights for hidden -> output
if self.regression_problem:
    self.b_out = np.ones(n_output) * np.mean(y_train)
else:
    self.b_out = np.zeros(n_output)
if self.initializer == 'Xavier':
    self.w_out = self.random.normal(loc=0.0, \
                                     scale=(2 / (self.n_hidden + n_output)
                                             , size=(self.n_hidden, n_output))
else:
    self.w_out = self.random.normal(loc=0.0, scale=0.1, \
                                     size=(self.n_hidden, n_output))

epoch_strlen = len(str(self.epochs)) # for progr. format.
self.eval_ = {'cost': [], 'train_acc': [], 'valid_acc': []}
if not self.regression_problem:
    y_train_enc = self._onehot(y_train, n_output)
else:
    y_train_enc = y_train
# iterate over training epochs
for i in range(self.epochs):
    # iterate over minibatches
    indices = np.arange(X_train.shape[0])
    if self.shuffle:
        self.random.shuffle(indices)
    for start_idx in range(0, \
                           indices.shape[0] - self.minibatch_size + 1, \
                           self.minibatch_size):
        batch_idx = indices\
        [start_idx:start_idx+self.minibatch_size]
        # forward propagation
        # z_h: list(z_h_1, z_h_2, ..... )
        z_h, a_h, z_out, a_out = \
        self._forward(X_train[batch_idx])
        #####
        # Backpropagation
        #####
        # [n_samples, n_classlabels]
        sigma_out = a_out - y_train_enc[batch_idx]

    for l in range(self.n_layers - 2):
        # [n_samples, n_hidden]
        # a_h: list(a_h_1, a_h_2, ..... )
        sigmoid_derivative_h = a_h[-1-1] * (1. - a_h[-1-1])
        # [n_samples, n_classlabels] dot
        # [n_classlabels, # n_hidden]
        # -> [n_samples, n_hidden]
        if l == 0:
            if self.regression_problem:

```

```

#         self.k['sigma_h_{0}'.format(self.n_layers-2-l)] = \
#         (np.dot(sigma_out, self.w_out.T))
#     else:
#         self.k['sigma_h_{0}'.format(self.n_layers-2-l)] = \
#         (np.dot(sigma_out, self.w_out.T) * sigmoid_derivative_h)
#     if self.is_sigmoid:
#         self.k['sigma_h_{0}'.format(self.n_layers-2-l)] = \
#         (np.dot(sigma_out, self.w_out.T) * sigmoid_derivative_h)
#     else:
#         self.k['sigma_h_{0}'.format(self.n_layers-2-l)] = \
#         (np.dot(sigma_out, self.w_out.T))
#     else:
#         if self.is_sigmoid:
#             self.k['sigma_h_{0}'.format(self.n_layers-2-l)] = \
#             (np.dot(self.k['sigma_h_{0}'.format(self.n_layers-1-l)]
#                     self.k['w_h_{0}'.format(self.n_layers-1-l)].T)
#             sigmoid_derivative_h)
#         else:
#             self.k['sigma_h_{0}'.format(self.n_layers-2-l)] = \
#             (np.dot(self.k['sigma_h_{0}'.format(self.n_layers-1-l)]
#                     self.k['w_h_{0}'.format(self.n_layers-1-l)].T)

for l in range(2, self.n_layers - 1):
    # [n_features, n_samples] dot [n_samples, n_hidden]
    # -> [n_features, n_hidden]
    self.k['grad_w_h_{0}'.format(l)] = np.dot(a_h[l-2].T, self.k['sigma_h_{0}'.format(l)])
    self.k['grad_b_h_{0}'.format(l)] = np.sum(self.k['sigma_h_{0}'.format(l)], axis=0)
    # [n_features, n_samples] dot [n_samples, n_hidden]
    # -> [n_features, n_hidden]
    self.k['grad_w_h_1'] = np.dot(X_train[batch_idx].T, self.k['sigma_h_1'])
    self.k['grad_b_h_1'] = np.sum(self.k['sigma_h_1'], axis=0)
    # [n_hidden, n_samples] dot [n_samples, n_classlabels]
    # -> [n_hidden, n_classlabels]
    grad_w_out = np.dot(a_h[-1].T, sigma_out)
    grad_b_out = np.sum(sigma_out, axis=0)

    # Regularization and weight updates
    for l in range(self.n_layers - 2):
        self.k['delta_w_h_{0}'.format(l+1)] = (self.k['grad_w_h_{0}'.format(l+1)] +
                                                self.l2*self.k['w_h_{0}'.format(l+1)])
        self.k['delta_b_h_{0}'.format(l+1)] = self.k['grad_b_h_{0}'.format(l+1)]
        self.k['w_h_{0}'.format(l+1)] -= self.eta * self.k['delta_w_h_{0}'.format(l+1)]
        self.k['b_h_{0}'.format(l+1)] -= self.eta * self.k['delta_b_h_{0}'.format(l+1)]
    delta_w_out = (grad_w_out + self.l2*self.w_out)
    delta_b_out = grad_b_out # bias is not regularized
    self.w_out -= self.eta * delta_w_out
    self.b_out -= self.eta * delta_b_out

#####
# Evaluation
#####
# Evaluation after each epoch during training
z_h, a_h, z_out, a_out = self._forward(X_train)
cost = self._compute_cost(y_train_enc, a_out)
y_train_pred = self.predict(X_train)
y_valid_pred = self.predict(X_valid)

```

```

        if not self.regression_problem:
            train_acc = ((np.sum(y_train == \
                                y_train_pred)).astype(np.float) /
                           X_train.shape[0])
            valid_acc = ((np.sum(y_valid == \
                                y_valid_pred)).astype(np.float) /
                           X_valid.shape[0])
        else:
            train_acc = self.R_square(output=y_train_pred, targets=y_train)
            valid_acc = self.R_square(output=y_valid_pred, targets=y_valid)
        self.eval_['cost'].append(cost)
        self.eval_['train_acc'].append(train_acc)
        self.eval_['valid_acc'].append(valid_acc)
#     return self
def get_weights(self):
    for i in range(self.n_layers - 2):
        print('w_h_{}: {}'.format(i + 1, self.k['w_h_{}'.format(i + 1)].shape)
              self.k['w_h_{}'.format(i + 1)])
    print('w_out: {}'.format(self.w_out.shape), self.w_out)
def get_loss_plot(self):
    plt.plot(range(self.epochs), self.eval_['cost'])
    plt.ylabel('Cost')
    plt.xlabel('Epochs')
    plt.show()

```

```

In [2]: import warnings
warnings.filterwarnings('ignore')
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.metrics import r2_score

```

data

```
In [3]: from sklearn.datasets import load_boston
X, y = load_boston(return_X_y=True)
boston = load_boston()
print(boston.DESCR)
```

```
.. _boston_dataset:
```

```
Boston house prices dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 506
```

```
:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.
```

```
:Attribute Information (in order):
```

```
  - CRIM      per capita crime rate by town
  - ZN        proportion of residential land zoned for lots over 25,000 sq.ft.
  - INDUS     proportion of non-retail business acres per town
  - CHAS      Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
  - NOX       nitric oxides concentration (parts per 10 million)
  - RM        average number of rooms per dwelling
  - AGE       proportion of owner-occupied units built prior to 1940
  - DIS       weighted distances to five Boston employment centres
  - RAD       index of accessibility to radial highways
  - TAX       full-value property-tax rate per $10,000
  - PTRATIO   pupil-teacher ratio by town
  - B         1000(Bk - 0.63)^2 where Bk is the proportion of blacks by town
  - LSTAT     % lower status of the population
  - MEDV      Median value of owner-occupied homes in $1000's
```

```
:Missing Attribute Values: None
```

```
:Creator: Harrison, D. and Rubinfeld, D.L.
```

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/> (<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>)

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression

problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

```
In [4]: df = pd.DataFrame(X, columns=boston.feature_names)
df.head(2)
```

Out[4]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.9	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.9	9.14

```
In [5]: df.describe()
```

Out[5]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.00
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.79
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.10
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.12
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.10
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.20
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.18
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.12

```
In [6]: f_dummy = np.array(df['CHAS'])[:, np.newaxis]
del df['CHAS']
f = np.array(df)
```

```
In [7]: from sklearn.preprocessing import StandardScaler
scalar = StandardScaler()
f = scalar.fit_transform(f)

X = np.concatenate((f, f_dummy), axis=1)
X.shape
```

Out[7]: (506, 13)

```
In [8]: split_idx = int(len(X) * 0.9)
y = y[:, np.newaxis]
X_train, X_test = X[:split_idx], X[split_idx:]
y_train, y_test = y[:split_idx], y[split_idx:]
```


train

```
In [9]: idx = int(len(X_train) * 0.9)
nn = NeuralNetMLP(epochs=50, dropout_rate=0.5, n_hidden=5, n_layers=3, minibatch_size=10)
nn.fit(X_train=X_train[:idx], \
      y_train=y_train[:idx], \
      X_valid=X_train[idx:], \
      y_valid=y_train[idx:])
```

```
In [10]: nn.get_weights()
```

```
w_h_1: (13, 5) [[ 0.91465083  1.02677732 -0.9625498  -0.86549987 -1.06891682]
 [-1.21793391 -1.08328517  0.95387523  1.35420732  0.98387599]
 [ 1.17331028  1.08411514 -1.17962309 -1.29020613 -1.33110834]
 [ 0.497263    0.74107993 -1.10802003 -0.90549064 -1.04737985]
 [-8.7014977  -6.9412538   6.40470926  7.88121075  6.3499206 ]
 [ 0.17936506  0.64279604 -0.81606997 -0.37955351 -0.79747339]
 [ 2.27258802  1.59794487 -1.09993141 -1.95366368 -1.00514567]
 [-1.12869033 -0.57984266  0.24115697  0.87340431  0.36835915]
 [ 0.25562907  0.72781903 -0.74809453 -0.46790867 -0.79915811]
 [ 2.40805552  2.42575578 -2.17382311 -2.6261473  -1.94778534]
 [-0.16453019 -0.13731082  0.39379391  0.17459239  0.49863058]
 [ 5.78599162  4.96261853 -4.93909243 -5.50816472 -4.71470803]
 [-0.89176109 -0.1819256   0.49315658  0.43121306  0.5519309  ]]
w_out: (5, 1) [[-3.31570253]
 [-2.58439414]
 [ 2.41904997]
 [ 2.94974384]
 [ 2.42346269]]
```

In [11]: `nn.get_loss_plot()`



```
In [12]: kk = locals()
n_models = 3
def nns(X_train, y_train, fig, axs, is_sigmoid=True, dropout_rate=0.8, n_layers=3):
    if is_sigmoid:
        kk['nn_0'] = NeuralNetMLP(n_layers=n_layers)
        kk['nn_1'] = NeuralNetMLP(initializer='Xavier', n_layers=n_layers)
        kk['nn_2'] = NeuralNetMLP(dropout_rate=dropout_rate, n_layers=n_layers)
    else:
        kk['nn_0'] = NeuralNetMLP(is_sigmoid=False, n_layers=n_layers)
        kk['nn_1'] = NeuralNetMLP(is_sigmoid=False, initializer='Xavier', n_layers=n_layers)
        kk['nn_2'] = NeuralNetMLP(is_sigmoid=False, dropout_rate=dropout_rate, n_layers=n_layers)
    for i in range(n_models):
        kk['nn_{}'.format(i)].fit(X_train=X_train[:idx], \
            y_train=y_train[:idx], \
            X_valid=X_train[idx:], \
            y_valid=y_train[idx:])

    title = ['default', 'Xavier', 'dropout']

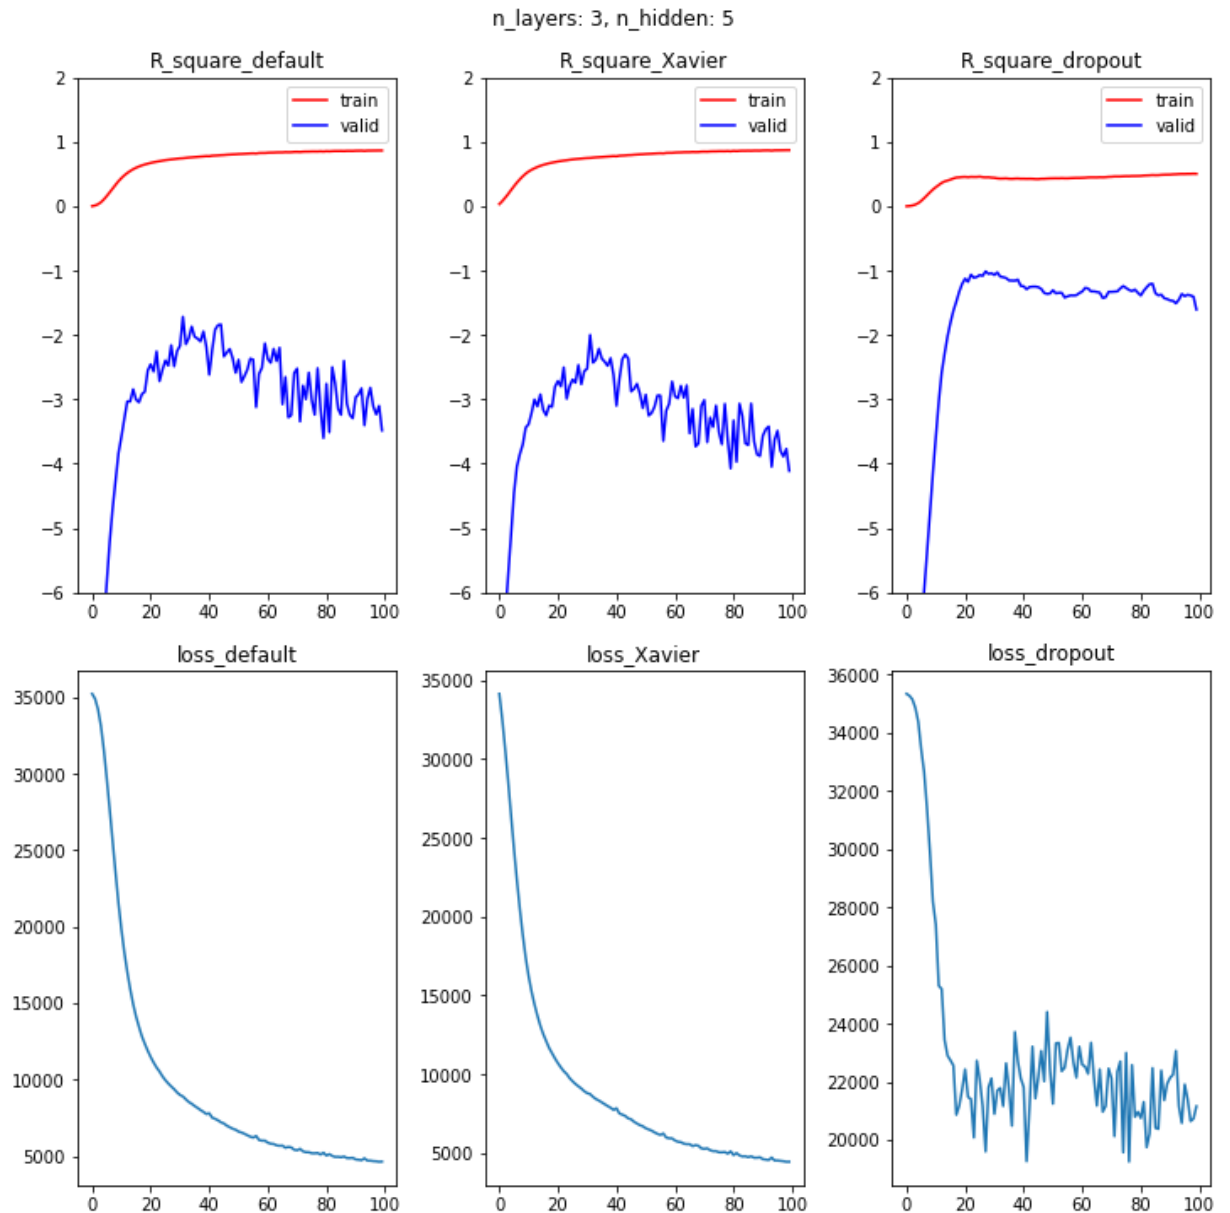
    col = ['red', 'blue']
    label = ['train', 'valid']

    for i in range(n_models):
        acc = [kk['nn_{}'.format(i)].eval_['train_acc'], \
            kk['nn_{}'.format(i)].eval_['valid_acc']]
        axs[0, i].set_title('R_square_' + title[i])
        axs[0, i].set_ylim([-6, 2])
        for l in range(len(col)):
            axs[0, i].plot(acc[l], color=col[l], label=label[l])
            axs[0, i].legend(frameon=True)
        axs[1, i].set_title('loss_' + title[i])
        axs[1, i].plot(range(kk['nn_{}'.format(i)].epochs), \
            kk['nn_{}'.format(i)].eval_['cost'])
    fig.suptitle('n_layers: {}, n_hidden: {}'.format(n_layers, nn_0.n_hidden))
    return fig
```

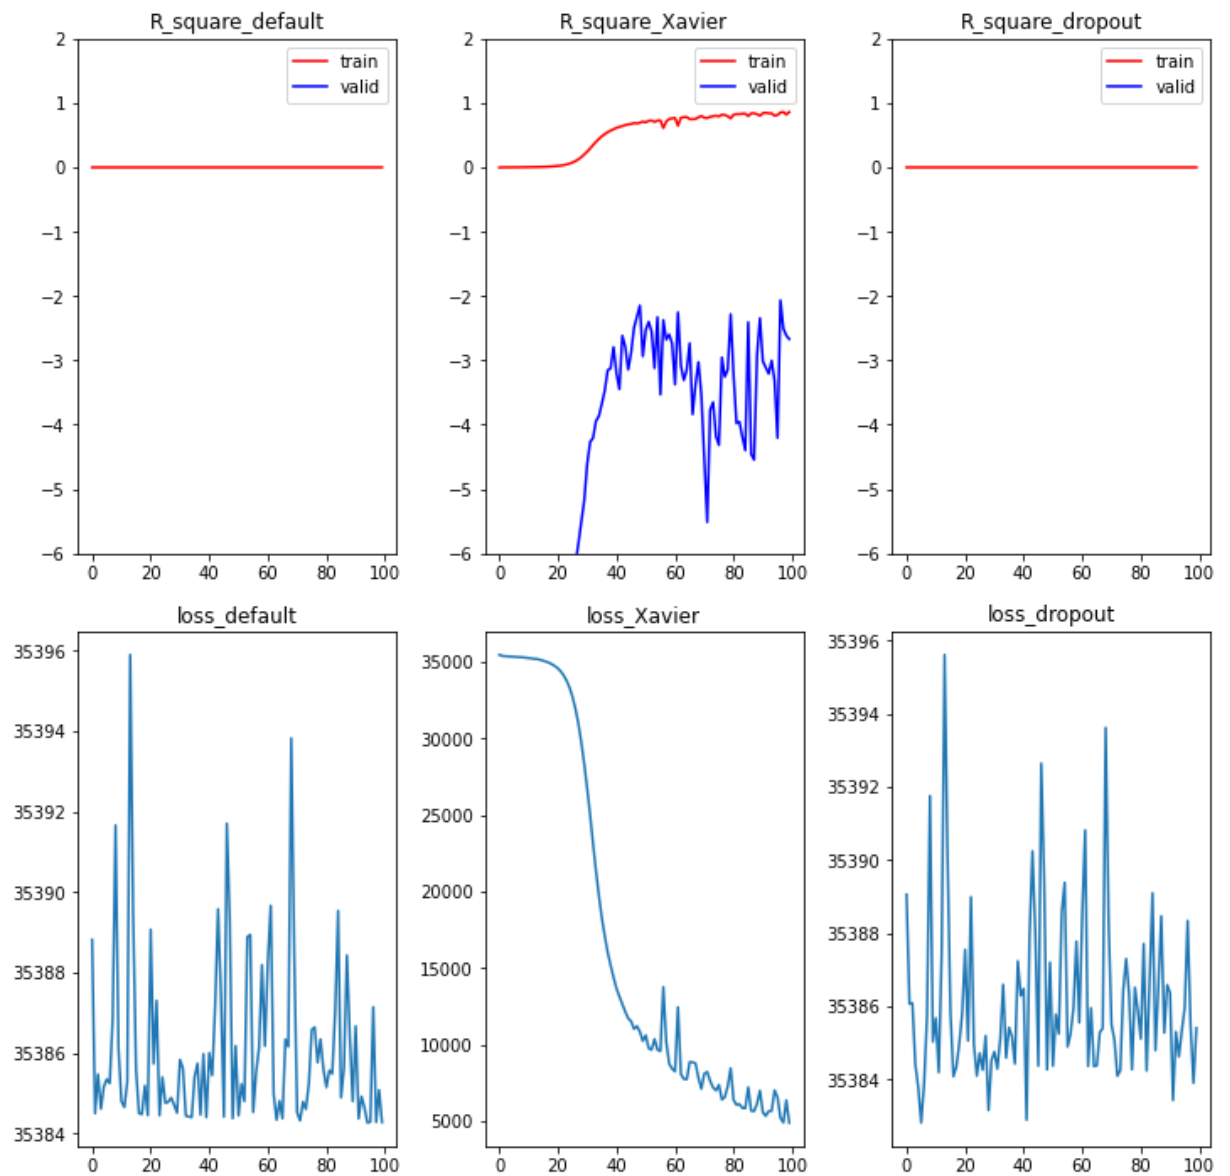
```

In [13]: n_n_layers = 3
for i in range(n_n_layers):
    n_layers = 3 + i*2
    kk['fig_{}'.format(i)], kk['axs_{}'.format(i)] = \
        plt.subplots(2, n_models, constrained_layout=True, figsize = (10, 10))
    kk['fig_{}'.format(i)] = nns(X_train, y_train, kk['fig_{}'.format(i)], kk['axs_{}'.format(i)],
                                is_sigmoid=True, dropout_rate=0.5, n_layers=n_layers)

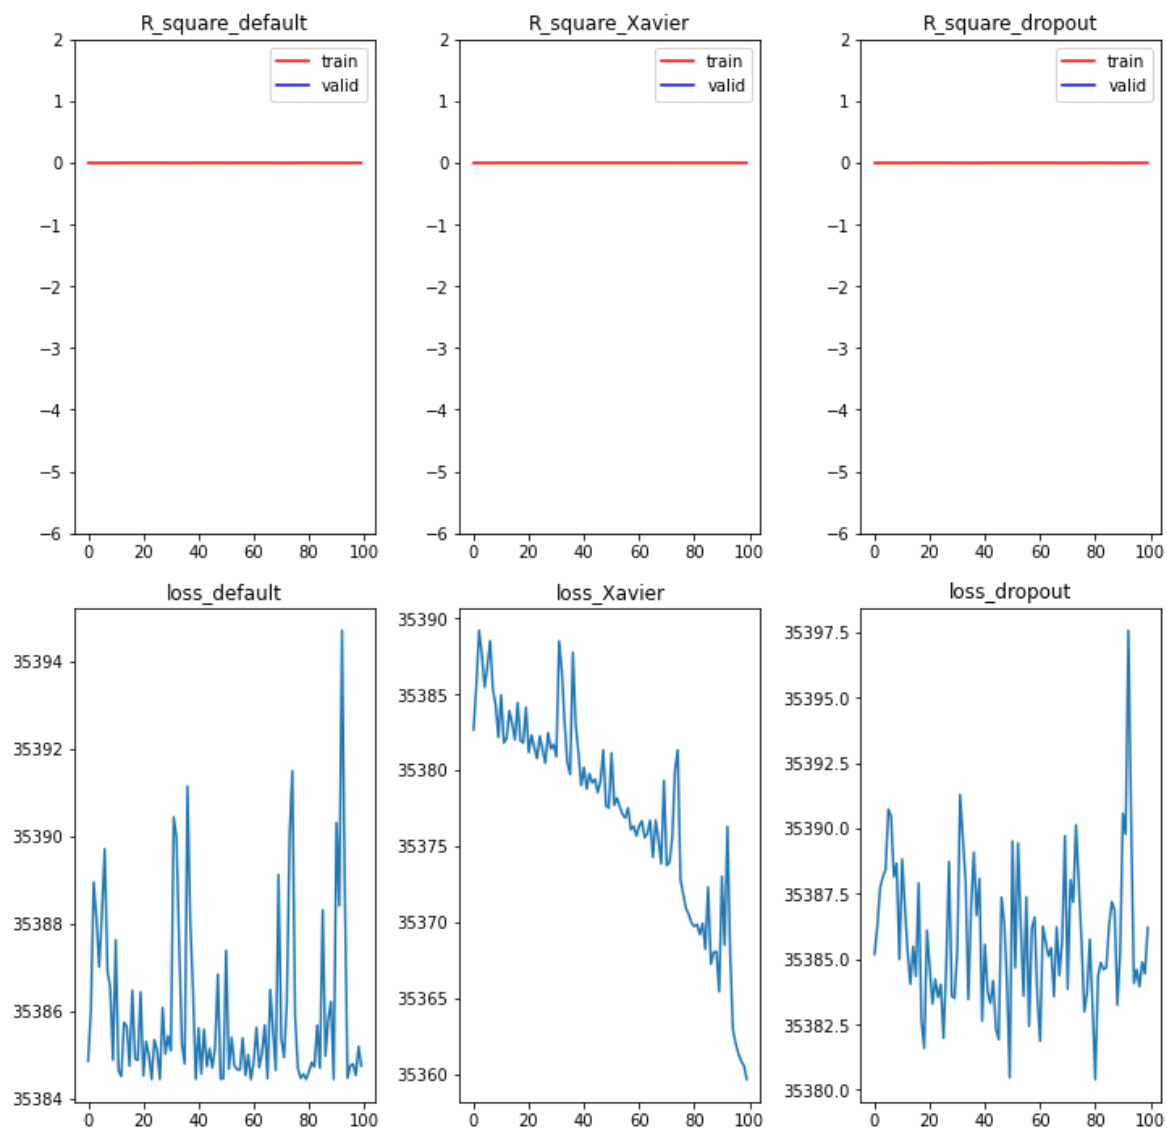
```



n_layers: 5, n_hidden: 5



n_layers: 7, n_hidden: 5



```
In [14]: y_fake = nn_0.predict(X_train) + np.random.normal(0, 1, y_train.shape)
```

```
In [15]: nn_0.predict(X_train)
```

```
[ 24.02514305],  
[ 24.02514302],  
[ 24.02514303],  
[ 24.02514304],  
[ 24.02514305],  
[ 24.02514305],  
[ 24.02514305],  
[ 24.02514305],  
[ 24.02514309],  
[ 24.0251431 ],  
[ 24.02514308],  
[ 24.02514307],  
[ 24.02514307],  
[ 24.02514306],  
[ 24.02514306],  
[ 24.02514305],  
[ 24.02514305],  
[ 24.02514301],  
[ 24.02514305],  
[ 24.02514306],
```

```
In [16]: y_fake
```

```
[ 23.61092158],  
[ 23.92484121],  
[ 24.04835094],  
[ 24.43245537],  
[ 24.74404654],  
[ 22.28661864],  
[ 23.30150156],  
[ 23.90328032],  
[ 23.64474879],  
[ 22.88075889],  
[ 24.48316083],  
[ 25.86448235],  
[ 23.48361724],  
[ 23.90861607],  
  
[ 25.91920143],  
[ 24.75466232],  
[ 24.30764747],  
[ 24.05531628],  
[ 25.84156892],  
[ 25.84156892],
```

```
In [17]: n_n_layers = 3
for i in range(n_n_layers):
    n_layers = 3 + i*2
    kk['fig_{}'.format(i)], kk['axs_{}'.format(i)] = \
        plt.subplots(2, n_models, constrained_layout=True, figsize = (10, 10))
    kk['fig_{}'.format(i)] = nns(X_train, y_fake, kk['fig_{}'.format(i)], kk['axs_{}'
        , is_sigmoid=False, dropout_rate=0.5, n_layers=n_layers)
```

